

IT UNIVERSITY IN COPENHAGEN

DEVOPS 2023

DevOps dudes: MiniTwit

Nikolaj Sørensen (nikso@itu.dk) Daniel Kjellberg (dakj@itu.dk)
Jeppe Lindhard (jepli@itu.dk) Johan Flensmark (jokf@itu.dk)
Benjamin Thygesen (beth@itu.dk)

Course Code: KSDSESM1KU
Course Manager: Helge Pfeiffer

24. May 2023

Contents

1	Introduction	1
2	System's Perspective	1
2.1	Architecture & Design	1
2.2	Dependencies	2
2.3	Current state of the system	4
2.4	License compatibility	5
3	Process's Perspective	5
3.1	The team	5
3.2	Version control setup	5
3.3	CI/CD	6
3.4	Monitoring & logging	6
3.5	Security assessment	8
3.6	Scaling & redundancy	10
3.7	AI assistance	10
4	Lessons Learned Perspective	10
4.1	DevOps style of work	11

1 Introduction

This report documents the creation of the application named *MiniTwit* which is made by the group named *dudes* for the DevOps course at IT-University of Copenhagen. The project is hosted by GitHub at github.com/Lindharden/DevOps, and the application is running at 157.230.76.157:8080.

In section 2 we describe the system and architecture of the MiniTwit application. Here we go through the dependencies of the application and how they interact with each other. In section 3 we describe the process of creating MiniTwit. This includes the team organization, and the setups that we used in order to do DevOps. Finally in section 4 we describe the issues we ran into underway, and the main lessons we have learned.

2 System's Perspective

This section documents the architecture of the MiniTwit application. We will describe the individual dependencies of the application and why we choose them. We will also describe what the current state of the system is.

2.1 Architecture & Design

The MiniTwit application relies on a lot of dependencies, which interact with each other in many ways. Figure 1 shows all of these interactions. In figure 1 it can be seen that a pipeline of actions is initiated every time someone commits to the MiniTwit GitHub repository. When someone commits a change, the GitHub Actions initialize our CI/CD setup, which is described in section 3.3. This setup deploys the image of the application and all its dependencies, which are described in section 2.2, to the Docker Hub. From the Docker Hub, everything is deployed to DigitalOcean which hosts our application. On DigitalOcean we use a replication setup, described in section 3.6, which ensures availability of the application.

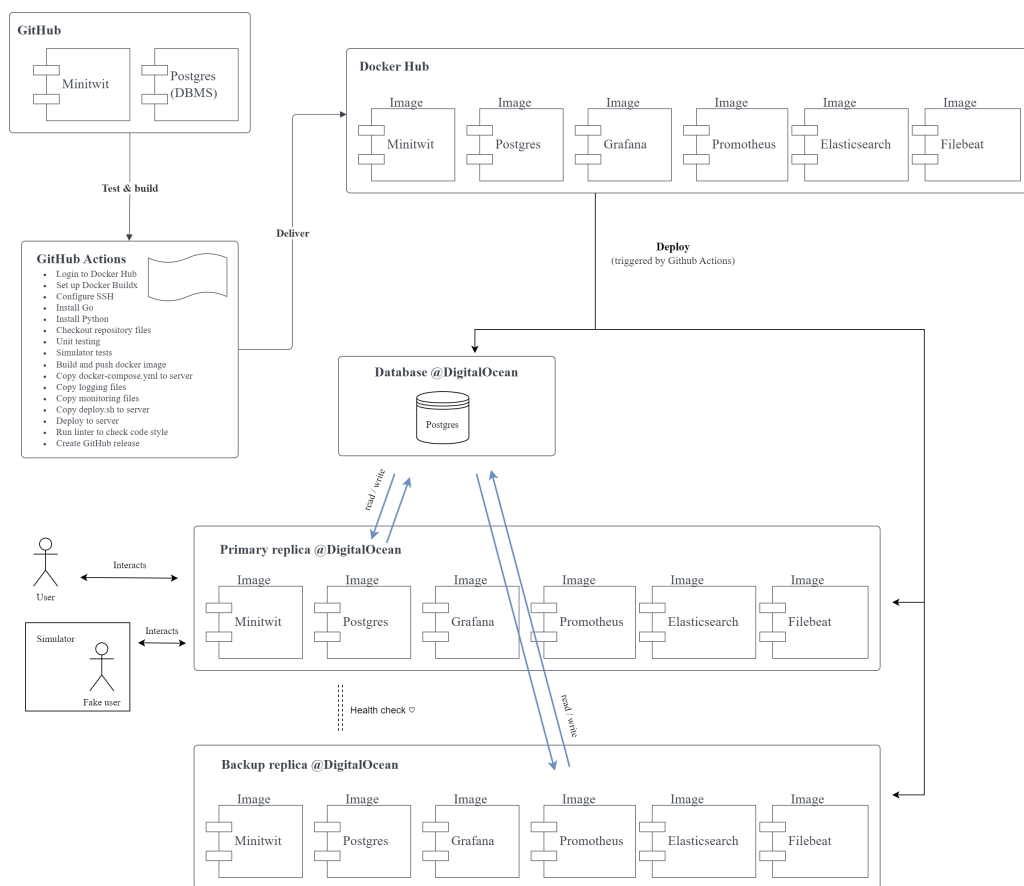


Figure 1: Architecture of the MiniTwit application. The diagram shows which dependencies are involved in the application, and how they interact with each other.

2.2 Dependencies

This section describes the different dependencies used in our application, and we argue for the choice of each technology.

Docker We use Docker¹ to containerize all of our dependencies, into convenient images which are far easier to transport and deploy. Without utilizing containerization, like Docker offers, it would be a struggle to keep track of every tool we use and all the technologies their depend on, and it would require a lot of steps every time we would have to run our application. Docker allows us to package everything together, and lets us run everything from one single command. The reason for using Docker, over other containerization tools such as Kubernetes, is because we already have experience using Docker from previous courses, and because it offers all the functionality we need for our application.

¹[docker.com](https://www.docker.com)

Go & Gin We choose to utilize Gin² for our web-application. Gin is an web framework written in Go/Golang, which allows for development of fast and stable web services. We choose to use Gin as it's faster than many other web development frameworks, including Flask which is the original framework that the MiniTwit application was using. We wanted our web framework to support HTML rendering, and to support the use of relational databases, as we wanted to be able to reuse as much of the original MiniTwit application as possible. Gin supports both of these features while being scalable and offering type safety, which is important for writing maintainable code, and therefore we thought that Gin was very suitable for our MiniTwit application. Additionally we also saw this project as an perfect opportunity to learn Go, and to use Gin.

Gorm In Go, one of the most popular ORM libraries is Gorm³. We chose this library because it is one of the most mature libraries we could choose out of the possible ones for Go. Gorm has the basic ORM features we needed like creation, update, and deletion of database objects. With Gorm we could create objects in Go that has the fields needed for the database.

Vagrant We utilize Vagrant⁴ in our project to easily configure and deploy the virtual machine that is running the application, onto the server at DigitalOcean. Vagrant allows us to give an generic description of the virtual machine we need, which it uses to create the virtual machine. This alleviates us from having to use manual programs such as VirtualBox where we would have to click around a UI. By automating this process we also alleviate the risk of making a mistake every time we setup up our VM, as everything is done the same way every time. Another reason for picking Vagrant over other alternatives, is because Vagrant is very flexible and can utilize a lot of different containerization platforms, and integrates well with different development ecosystems. Additionally it also works on all the common operating systems such as Windows, macOS and Linux.

Prometheus & Grafana We use Prometheus⁵ to gain valuable insight about our Minitwit application for the following reasons: Prometheus supports a wide range of metrics. Prometheus works well with the Go programming language. Prometheus can easily be integrated with dashboard tools such as Grafana⁶. Prometheus can alert the team in the event that any metrics surpasses certain thresholds. Furthermore, Grafana simply visualizes the metrics to create a good overview of all the important metrics. We chose Grafana because it is easy to create a dashboard that only has to be configured once.

²gin-gonic.com

³gorm.io

⁴vagrantup.com

⁵prometheus.io

⁶grafana.com

Elasticsearch, Filebeat & Kibana To facilitate logging in our Minitwit application, we decided to implement an EFK stack. The EFK stack is a great fit for our Minitwit application as it provides a lightweight log shipper in Filebeat which can collect logs from the application, parse them and forward them to a storage component. It provides an efficient storage and analytics engine in Elastic search, which enables efficient handling of large data volumes and is excellent for full-text search. Lastly, it provides a simple to setup dashboard in Kibana that allows for filtering queries on the data stores in elastic search, such that we only display the data that we find most valuable.

DigitalOcean We have DigitalOcean⁷ as our cloud infrastructure provider. DigitalOcean offers a user-friendly interface for setting up and deploying Droplets (VMs) and Volumes (storage). In general, it is known for its simplicity and affordability which makes it ideal for our Minitwit application.

Github & GH Actions To facilitate our CI/CD workflow we utilize GitHub Actions because it is integrated in GitHub, which is the version control system we use for our repository. This way we don't have to involve other platforms, and we can keep everything in one place. Also, since it is integrated in GitHub, we can see the status of individual commits when we push them. This way we can see whether individual commits build (or contain errors), and we can see whether individual commits pass tests. Lastly, it allows for Continuous Delivery and Deployment, which means we can automate the delivery or deployment of our changes all with the same system.

2.3 Current state of the system

At the time of writing (20-05-2023) our static code analysis tools report the following:

- **Sonarcloud:** Latest pull request passed with 0 bugs, 0 vulnerabilities, 0 security hotspots and 10 code smells. We have assessed the code smells and decided to ignore them as they are insignificant.
- **Snyk:** It is currently disabled as it caused issues that blocked our CI/CD chain. We have not prioritized to fix the underlying problem as we believe the other code analysis tools provide a satisfying coverage of potential security vulnerabilities.
- **CodeQL:** Latest deployment passed CodeQL in the CI/CD chain.
- **Dependabot:** No active pull requests for vulnerable dependencies.
- **Lint for Go:** Latest deployment passed Lint in the CI/CD chain.

⁷digitalocean.com

2.4 License compatibility

Our project is under GPL 3.0 which is compatible⁸ with all dependencies, namely: *MIT License*, *GNU Affero General Public License v3.0*, *Apache License 2.0*, *Server Side Public License*, *ISC*, *BSD-3-Clause*.

3 Process's Perspective

In this section we present how we organized the development process.

3.1 The team

Dudes is a *close-knit team* consisting of five Master students. We value an in-person working environment when possible, which also means that a majority of project work and general interactions have taken place in-person at school during the course hours. With our collective expertise, we have tackled most challenges collaboratively, working together to find solutions. This approach has fostered a seamless workflow, enabling efficient problem-solving and encouraging a sense of shared accomplishment within our team.

Our team has utilized GitHub Issues as a centralized tool to track and manage our tasks. By documenting each task as an issue, we have maintained a clear overview of our work and assigned responsibilities. We implemented issue templates in GitHub, providing a standardized structure for creating issues, ensuring that all essential details were included consistently. This approach enhanced clarity and efficiency in our communication, making it easier for team members to understand and address the tasks at hand.

3.2 Version control setup

Our version control system setup revolves around a monorepo approach, providing a centralized repository to manage our codebase. We have chosen this approach to consolidate all related code in a single repository, promoting the ease of code reuse and streamlining collaboration. To enhance our workflow and issue tracking, we utilize GitHub's built-in issues extensively. The issues created in GitHub follows predefined templates, defined with GitHub issue templates, in order to ensure consistent issue management and to establish guidelines for issue creation, assignment, and resolution. These templates provide a structured framework for issue creation, ensuring all necessary information is captured for the respective type of issue, whether it be bug reports, feature requests, or security vulnerabilities. Additionally, we have implemented a well-defined branching strategy, consisting of a main branch as the primary branch, and subbranches designated for the development of single features. The subbranches are named according to a set of naming schemas in order to provide clarity and organization. The naming schemas consists of a predefined tag followed

⁸gnu.org/licenses/license-list.en.html

by a relevant name, such as "feature/feature-name" or "bug/bug-name". Before merging any changes into the main branch, a set of checks is performed to ensure code quality and stability. This branching strategy enables seamless feature development, bug fixing, and code maintenance, ensuring a streamlined development process and promoting collaboration within our team. To track releases and provide visibility into the version history, we utilize the GitHub Releases feature. This feature is configured to automatically create a new release for each deployment, ensuring the latest release is always up-to-date with the deployed version. In addition to the automatic releases, we also implemented a weekly release schedule. This manual process allows us to consolidate any significant and relevant features or updates that have been developed during the week. By creating weekly releases, we provide a clear and structured approach to showcasing our progress and informing about the latest developments.

3.3 CI/CD

For our CI/CD setup, we leveraged the capabilities of GitHub Actions to automate the verification process when code is merged into the main branch. GitHub Actions allows us to define custom workflows that encompasses various stages of our pipeline. Within these workflows, we incorporated essentials steps to ensure code quality and reliability.

Firstly, the actions execute unit tests to validate the functionality of our codebase, ensuring that it meets the expected behavior. Secondly, we integrated static analysis tools into the workflow, specifically Lint for Go, Snyk, CodeQL, and SonarCloud. These third-party services thoroughly analyze our code to detect potential security vulnerabilities, code smells, and many other potential issues, assisting in maintaining a high level of code quality. Lastly, once all the checks pass successfully, GitHub automatically builds a Docker image of our project. This image is then deployed to our hosting service, ensuring that the latest version of our application is readily available to our users.

To further enhance our development workflow, we integrated GitHub Actions with the pull request feature. This integration enables us to run tests and analysis on proposed changes before merging them into the main branch, minimizing the risk of introducing faulty code into the main branch.

The integration capabilities and the support for automated tests provided by GitHub Actions were key factors in our decisions to choose this CI/CD solution. Our setup streamlines the development process, automates quality checks, and ensures our code is validated and ready for deployment upon introducing changes.

3.4 Monitoring & logging

The monitoring system we have implemented for Minitwit is Prometheus. We have implemented application monitoring and infrastructure monitoring through the following metrics:

- **Number of tweet requests per hour** - This helped us determine the amount of load our application is under. We will be able to see when something critical is happening to the application, which is arguably the most important feature.
- **Error codes** - This metric shows us the total amount of errors which the system reports. This metric can be used to determine whether everything runs as expected. If the dashboard displays an unusual amount of errors, this might be a hint that something is going on with the system.
- **Response time** - This metric tracks how many milliseconds a user has to wait for for a request response. This is useful to see if any network issues arise and also to ensure that we uphold the expected average response time of <50ms included in our API SLA.
- **Hardware load** - This metric reports system load, CPU and memory usage, disk I/O and amount of network transmit and receive requests. This can be useful for determining whether everything runs as expected and whether the hosting server has sufficient resources to handle the incoming load.

We use a Grafana dashboard to display the data collected for each metric through Prometheus, which can be seen in figure 2. The dashboard is configured through a JSON file.⁹ Both Prometheus and Grafana have been integrated into the application through the *docker-compose.yml* configuration file.

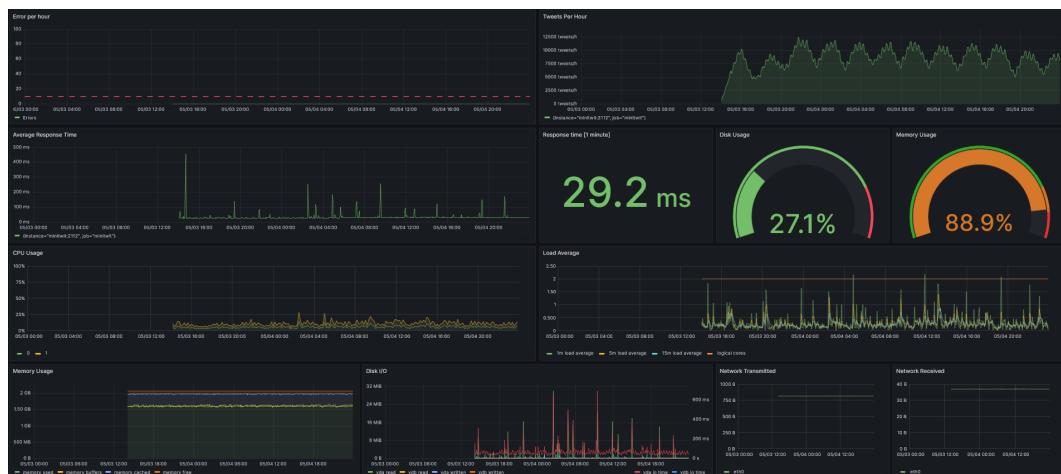


Figure 2: Grafana dashboard displaying the data collected with Prometheus.

The logging system we have implemented an EFK stack consisting of Elastic search, Filebeat and Kibana. We used the *Zap* package for Go to nicely format and emit log messages that Filebeat collects. We have two types of log messages, the first being the native Gin Gonic

⁹ github.com/Lindharden/DevOps/blob/main/monitoring/grafana/dashboards/main.json.

log messages that happen for every request. The second one is the log messages that we have implemented in the code, which are formatted by *Zap* in JSON as follows:

```
{
  "level": "error",
  "timestamp": "2023-03-25T13:09:38Z",
  "caller": "controllers/loginController.go:130",
  "msg": "Could not find session",
  "user": "John Doe",
  "stacktrace": "..."}
}
```

The Kibana dashboard can be seen in figure 3. The EFK stack is integrated into the application through the docker container.

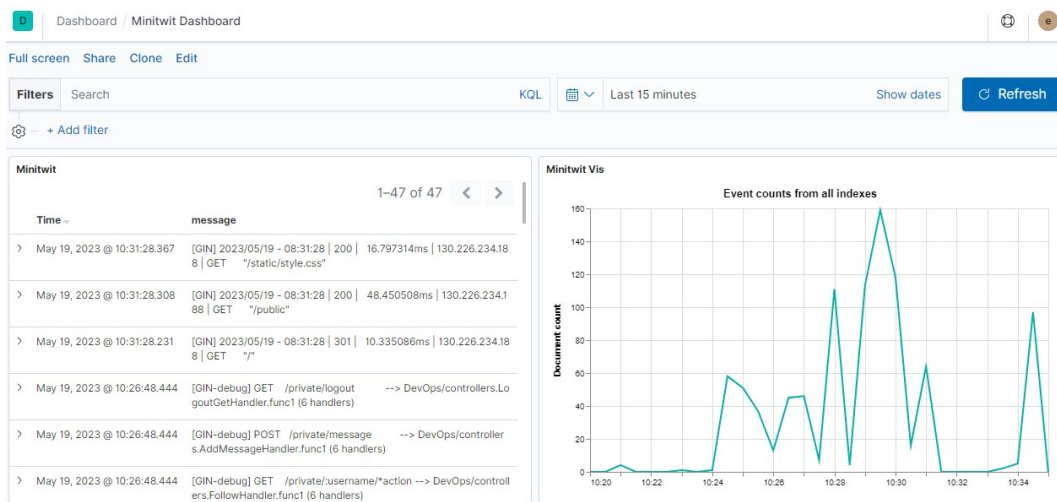


Figure 3: Kibana dashboard displaying the logs that are collected and organised through Filebeat and Elastic search.

3.5 Security assessment

In our security assessment we have considered which risks the infrastructure and application are vulnerable to. The risks we have assessed are mapped in a risk assessment matrix which can be seen in figure 4. More concretely they are:

- **Compromised Github repository** - The github repository owner uses two-factor authentication. Furthermore, all users must make a pull request and have it approved to change the code on the main branch. This makes it unlikely that the github repository is compromised. However, a compromised GitHub repository would mean that the attacker would gain access to one of our GitHub admin accounts, which would mean they have the ability to change our code, making it a significant impact risk.

- **Whitebox exploits** - As our GitHub repository is public, anyone can see our code and create exploits for it. Therefore if the attacker finds a reason and an opportunity to create a whitebox exploit, it would likely happen. However, as we continuously update our code and update vulnerable packages the chance of an exploit being possible, is very low. As the attacker knows what our code base looks like they can plan their attacks towards having as large of an impact as possible. In addition, the attacker can see all the dependencies we use, meaning they can utilize any exploits found in our dependencies, making it an extensive impact risk.
- **DDoS attack** - Since our application is not that popular it is not likely that it will get DDoS'ed, but it is still possible since our application is public. A DDoS attack would mean that our application would be unreachable for a period of time. This would be a moderate inconvenience for us and our users, making it a moderate impact risk.
- **SQL injection** - Since the database abstraction layer library GORM sanitizes all keywords, an SQL injection attack is very rare, basically impossible. SQL injection would mean that an attacker performs attacks directly on our database, by sending arbitrary SQL queries to our database, making it an extensive impact risk.
- **DB leak** - The only ways our database can be leaked, is by someone compromising our code repository and performing some kind of whitebox exploit, or by performing SQL injection. As we deemed these risks as "rare", our database leaking must also be a rare event. If our database leaked, the attacker would gain access to usernames, email's and hashed passwords of our users. As none of this is sensitive information the impact would be negligible. However, the hashed passwords could be used in an brute-force attempt to identify passwords, making it a negligible impact risk.




Figure 4: Risk assessment matrix. The vertical axis indicates the probability of a risk while the horizontal axis indicates the impact of a risk. Risks in the top-right corner are considered problematic and should be handled accordingly.

None of the assessed risks are placed inside the problematic area in the risk assessment matrix in figure 4. Therefore, we have not taken any actions against these risks, we have simply decided to accept them for now.

3.6 Scaling & redundancy

To ensure fault tolerance we introduced redundancy in the form of a replica. We created an additional droplet on DigitalOcean and changed our Vagrantfile accordingly to have identical instances running on each of the droplets. One of the instances act as the primary one, for which all traffic is directed to. The other instance serves as a backup. With the use of heartbeats the server can keep track of the primary replicas state, and should it fail to show signs of life, the backup replica will take over as primary.

In addition to fault tolerance, having a primary and backup replica also enables rolling updates which eliminates the inevitable downtime caused from deploying updates. We did not get around to do it for this project, but we had a plan of how we wanted to do it. In order to increase our availability, we will first be updating the backup replica and make it our primary server, then we can shut down the replica which was previously the primary one and update it as well. DigitalOcean will make this easy as it automatically switches the primary server to be the replica which is running.

3.7 AI assistance

For project work we have consulted ChatGPT for setting up Docker, however, without much benefit. For writing the report we have used ChatGPT in a couple of places to make some paragraphs more concise.

4 Lessons Learned Perspective

During the development of the project, we had some issues while setting up new services in Docker. We did not have prior experience with many of the services such as Grafana, Prometheus, and Kibana, and therefore we found that it took way longer than expected to get these services to run properly. There was not a simple solution, other than trying until we got it right, and this was partly because of how hard it was to narrow down problems as the Docker logs are not always that elaborate. From this we learned to dedicate more time to tasks that involve the implementation of new services.

Our project utilize many different tools for code quality, and security assesment, hereunder Snyk. We implemented Snyk into our CI/CD chain, where it's supposed to stop the deployment if it finds any major problems with the commit. This ran fine for a while until there suddenly went something wrong in the setup. From this point on Snyk failed our CI/CD pipeline, not because of any vulnerabilities, but because of some other unidentified

problems¹⁰. We didn't find any solution to this, so we solved the problem by disabling Snyk entirely. We deemed this fine for now, as we have other tools which also check for vulnerabilities in our application.

4.1 DevOps style of work

For a majority of other school related development projects we have not utilized a concrete tool to organise tasks and issues. For this project we took great advantage of GitHub Issues to have a centralized backlog that was also streamlined using issue templates. GitHub Issues provided a much needed overview of features and bugs and made it easy to manage and distribute responsibilities.

The DevOps handbook [Kim+21] mentions the Three Ways. This is our thoughts on how we adhered to each of the principles:

- **Flow:** We adhere to the Flow principle which talks about implementing a fast left-to-right flow, which is the time it takes from when requests are put on the backlog, till they are implemented and is in production. Our setup has a fast Flow as we have an CI/DC setup which automatically builds, tests and deploys changes from the main branch, to our remote server at DigitalOcean. This means that the time it takes from when changes are committed, till they are in production, is just as long as it takes for our team to approve and merge pull-requests. We also try to split tasks into as small as possible backlog items (GitHub Issues), such that they are easier to delegate to different group members, such that they are passed through the pipeline faster.
- **Feedback:** We adhere to the Feedback principle, which focuses on continuous problem solving when they occur. Our application don't have real users, but we simulate the feedback process by receiving error messages from the user simulation run by the teaching team. When we see that errors occur, we create an Issue/ticket which goes in our backlog. We then try to fix the error as fast as possible.
- **Continual Learning and Experimentation:** We try to adhere to Continual Learning and Experimentation where we encourage risk taking and see mistakes as opportunities to learn. We are in a special situation as we are not implementing a real application, and therefore we can afford to make big mistakes without any bigger impacts. We relieve the pressure of having to manually deploy the changes we make, by having an CI/CD setup which automatically deploys all the changes we make.

¹⁰Problem with Snyk can be seen in GitHub Actions: github.com/Lindharden/DevOps/actions/runs/4660817522.

Literature

- [Kim+21] Gene Kim et al. *The DevOps handbook: How to create world-class agility, reliability, & security in technology organizations*. IT Revolution, 2021. Chap. 1-4.